

STATS205P_Project

Chuqi Wang

2024-06-12

```
library(rstan)
```

```
## Loading required package: StanHeaders
```

```
##
```

```
## rstan version 2.32.6 (Stan version 2.32.2)
```

```
## For execution on a local, multicore CPU with excess RAM we recommend calling
```

```
## options(mc.cores = parallel::detectCores()).
```

```
## To avoid recompilation of unchanged Stan programs, we recommend calling
```

```
## rstan_options(auto_write = TRUE)
```

```
## For within-chain threading using `reduce_sum()` or `map_rect()` Stan functions,
```

```
## change `threads_per_chain` option:
```

```
## rstan_options(threads_per_chain = 1)
```

```
library(rstan)
```

```
getwd()
```

```
## [1] "/Users/chuqi wang/Desktop/UCI/STATS205P/Project"
```

```
data = read.csv("healthcare-dataset-stroke-data.csv", header = TRUE,  
               na.strings = c("N/A", "NA", ""))
```

```
summary(data)
```

```
##           id           gender           age           hypertension  
## Min.      : 67   Length:5110   Min.      : 0.08   Min.      :0.00000  
## 1st Qu.:17741   Class :character   1st Qu.:25.00   1st Qu.:0.00000  
## Median :36932   Mode  :character   Median :45.00   Median :0.00000  
## Mean    :36518                      Mean    :43.23   Mean    :0.09746  
## 3rd Qu.:54682                      3rd Qu.:61.00   3rd Qu.:0.00000  
## Max.    :72940                      Max.    :82.00   Max.    :1.00000  
##  
## heart_disease   ever_married   work_type   Residence_type  
## Min.      :0.00000   Length:5110   Length:5110   Length:5110  
## 1st Qu.:0.00000   Class :character   Class :character   Class :character  
## Median :0.00000   Mode  :character   Mode  :character   Mode  :character  
## Mean      :0.05401  
## 3rd Qu.:0.00000  
## Max.      :1.00000  
##  
## avg_glucose_level   bmi   smoking_status   stroke  
## Min.      : 55.12   Min.      :10.30   Length:5110   Min.      :0.00000  
## 1st Qu.: 77.25   1st Qu.:23.50   Class :character   1st Qu.:0.00000  
## Median : 91.89   Median :28.10   Mode  :character   Median :0.00000  
## Mean     :106.15   Mean      :28.89                      Mean     :0.04873
```

```
## 3rd Qu.:114.09    3rd Qu.:33.10    3rd Qu.:0.00000
## Max.    :271.74    Max.    :97.60    Max.    :1.00000
##                      NA's    :201
```

```
head(data)
```

```
##      id gender age hypertension heart_disease ever_married work_type
## 1  9046   Male  67           0             1         Yes   Private
## 2 51676 Female  61           0             0         Yes Self-employed
## 3 31112   Male  80           0             1         Yes   Private
## 4 60182 Female  49           0             0         Yes   Private
## 5  1665 Female  79           1             0         Yes Self-employed
## 6 56669   Male  81           0             0         Yes   Private
## Residence_type avg_glucose_level bmi smoking_status stroke
## 1      Urban      228.69 36.6 formerly smoked      1
## 2      Rural      202.21  NA   never smoked      1
## 3      Rural      105.92 32.5 never smoked      1
## 4      Urban      171.23 34.4      smokes      1
## 5      Rural      174.12 24.0 never smoked      1
## 6      Urban      186.21 29.0 formerly smoked      1
```

```
dim(data)
```

```
## [1] 5110  12
```

```
names(data)
```

```
## [1] "id"           "gender"       "age"
## [4] "hypertension" "heart_disease" "ever_married"
## [7] "work_type"    "Residence_type" "avg_glucose_level"
## [10] "bmi"         "smoking_status" "stroke"
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
data %>% summarise_all(~sum(is.na(.)))
```

```
##      id gender age hypertension heart_disease ever_married work_type
## 1  0      0  0           0             0         0         0
## Residence_type avg_glucose_level bmi smoking_status stroke
## 1      0           0 201           0         0
```

```
data$bmi[is.na(data$bmi)] <- median(data$bmi, na.rm = TRUE)
```

```
# Load necessary libraries
# Load necessary libraries
library(ggplot2)
library(gridExtra)
```

```
##
```

```

## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##      combine

library(grid)
# Create a density plot for age
p1 <- ggplot(data, aes(x = age)) +
  geom_density(fill = "blue", alpha = 0.5) +
  theme_minimal() +
  labs(title = "Density Plot of Age", x = "Age", y = "Density") +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Create a density plot for BMI
p2 <- ggplot(data, aes(x = bmi)) +
  geom_density(fill = "green", alpha = 0.5) +
  theme_minimal() +
  labs(title = "Density Plot of BMI", x = "BMI", y = "Density") +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Create a density plot for average glucose level
p3 <- ggplot(data, aes(x = avg_glucose_level)) +
  geom_density(fill = "purple", alpha = 0.5) +
  theme_minimal() +
  labs(title = "Density Plot of Average Glucose Level", x = "Average Glucose Level", y = "Density") +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Box plot for age by stroke status
p4 <- ggplot(data, aes(x = factor(stroke), y = age)) +
  geom_boxplot(fill = "blue", alpha = 0.5) +
  theme_minimal() +
  labs(title = 'Age by Stroke Status', x = 'Stroke', y = 'Age') +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Box plot for bmi by stroke status
p5 <- ggplot(data, aes(x = factor(stroke), y = bmi)) +
  geom_boxplot(fill = "green", alpha = 0.5) +
  theme_minimal() +
  labs(title = 'BMI by Stroke Status', x = 'Stroke', y = 'BMI') +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Box plot for avg_glucose_level by stroke status
p6 <- ggplot(data, aes(x = factor(stroke), y = avg_glucose_level)) +
  geom_boxplot(fill = "purple", alpha = 0.5) +

```

```

theme_minimal() +
labs(title = 'Average Glucose Level by Stroke Status', x = 'Stroke', y = 'Average Glucose Level') +
theme(plot.title = element_text(size = 15),
      axis.title = element_text(size = 12),
      axis.text = element_text(size = 10))

# Scatter plot for age vs stroke
p7 <- ggplot(data, aes(x = age, y = stroke)) +
  geom_jitter(color = "blue", alpha = 0.5) +
  theme_minimal() +
  labs(title = 'Age vs Stroke', x = 'Age', y = 'Stroke') +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

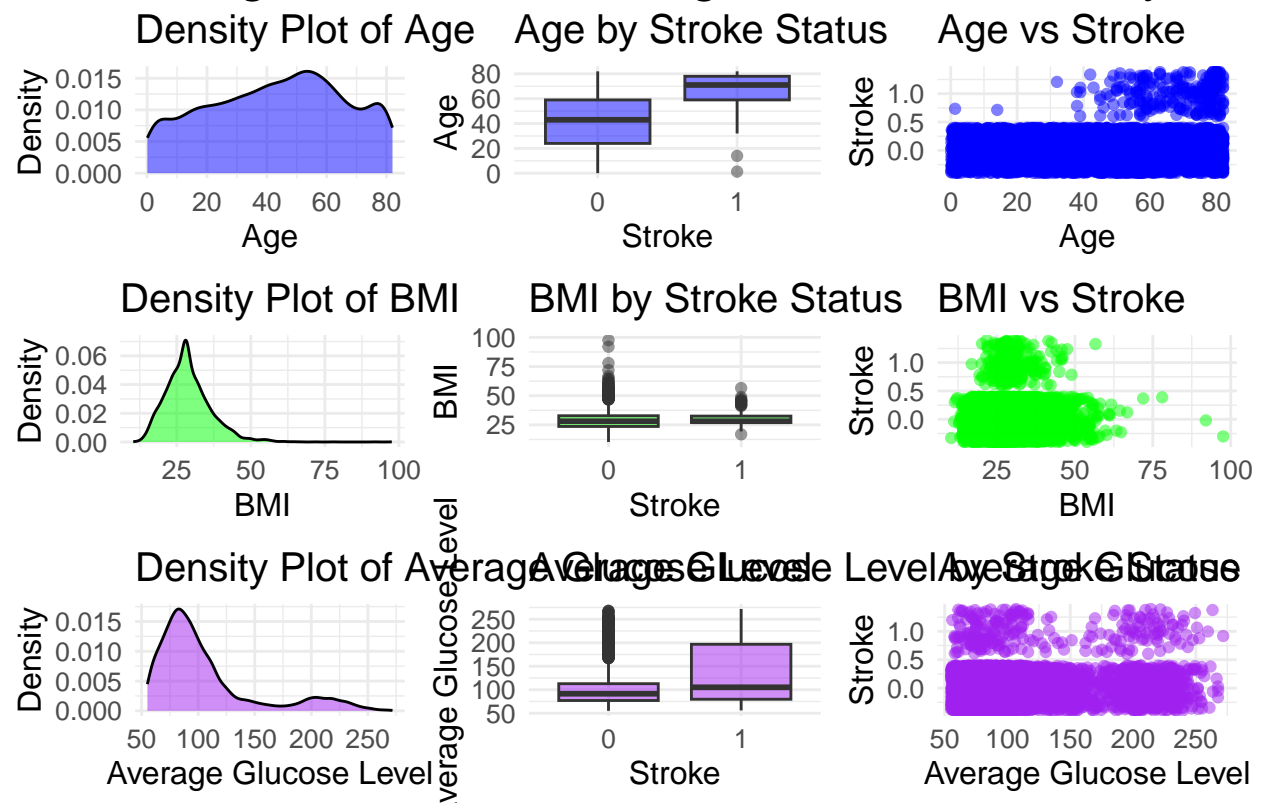
# Scatter plot for bmi vs stroke
p8 <- ggplot(data, aes(x = bmi, y = stroke)) +
  geom_jitter(color = "green", alpha = 0.5) +
  theme_minimal() +
  labs(title = 'BMI vs Stroke', x = 'BMI', y = 'Stroke') +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Scatter plot for avg_glucose_level vs stroke
p9 <- ggplot(data, aes(x = avg_glucose_level, y = stroke)) +
  geom_jitter(color = "purple", alpha = 0.5) +
  theme_minimal() +
  labs(title = 'Average Glucose Level vs Stroke', x = 'Average Glucose Level', y = 'Stroke') +
  theme(plot.title = element_text(size = 15),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

# Combine the three plots into one layout with increased figure size
grid.arrange(p1, p4, p7, p2, p5, p8, p3, p6, p9, ncol = 3, nrow = 3,
             top = textGrob("Combined Plots of Age, BMI, and Average Glucose Level by Stroke Status",
                           gp = gpar(fontsize = 20)))

```

Plots of Age, BMI, and Average Glucose Level by Str



```
# Create a bar plot for gender
p1 <- ggplot(data, aes(x = gender)) +
  geom_bar(aes(fill = gender), color = "black") +
  theme_minimal() +
  labs(title = "Gender", x = "Gender", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))

# Create a bar plot for hypertension
p2 <- ggplot(data, aes(x = factor(hypertension))) +
  geom_bar(aes(fill = factor(hypertension)), color = "black") +
  theme_minimal() +
  labs(title = "Hypertension", x = "Hypertension", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))

# Create a bar plot for heart_disease
p3 <- ggplot(data, aes(x = factor(heart_disease))) +
  geom_bar(aes(fill = factor(heart_disease)), color = "black") +
  theme_minimal() +
  labs(title = "Heart Disease", x = "Heart Disease", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
```

```

theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
      axis.title = element_text(size = 12),
      axis.text = element_text(size = 10, angle = 45, hjust = 1))

# Create a bar plot for ever_married
p4 <- ggplot(data, aes(x = ever_married)) +
  geom_bar(aes(fill = ever_married), color = "black") +
  theme_minimal() +
  labs(title = "Ever Married", x = "Ever Married", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))

# Create a bar plot for work_type
p5 <- ggplot(data, aes(x = work_type)) +
  geom_bar(aes(fill = work_type), color = "black") +
  theme_minimal() +
  labs(title = "Work Type", x = "Work Type", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))

# Create a bar plot for Residence_type
p6 <- ggplot(data, aes(x = Residence_type)) +
  geom_bar(aes(fill = Residence_type), color = "black") +
  theme_minimal() +
  labs(title = "Residence Type", x = "Residence Type", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))

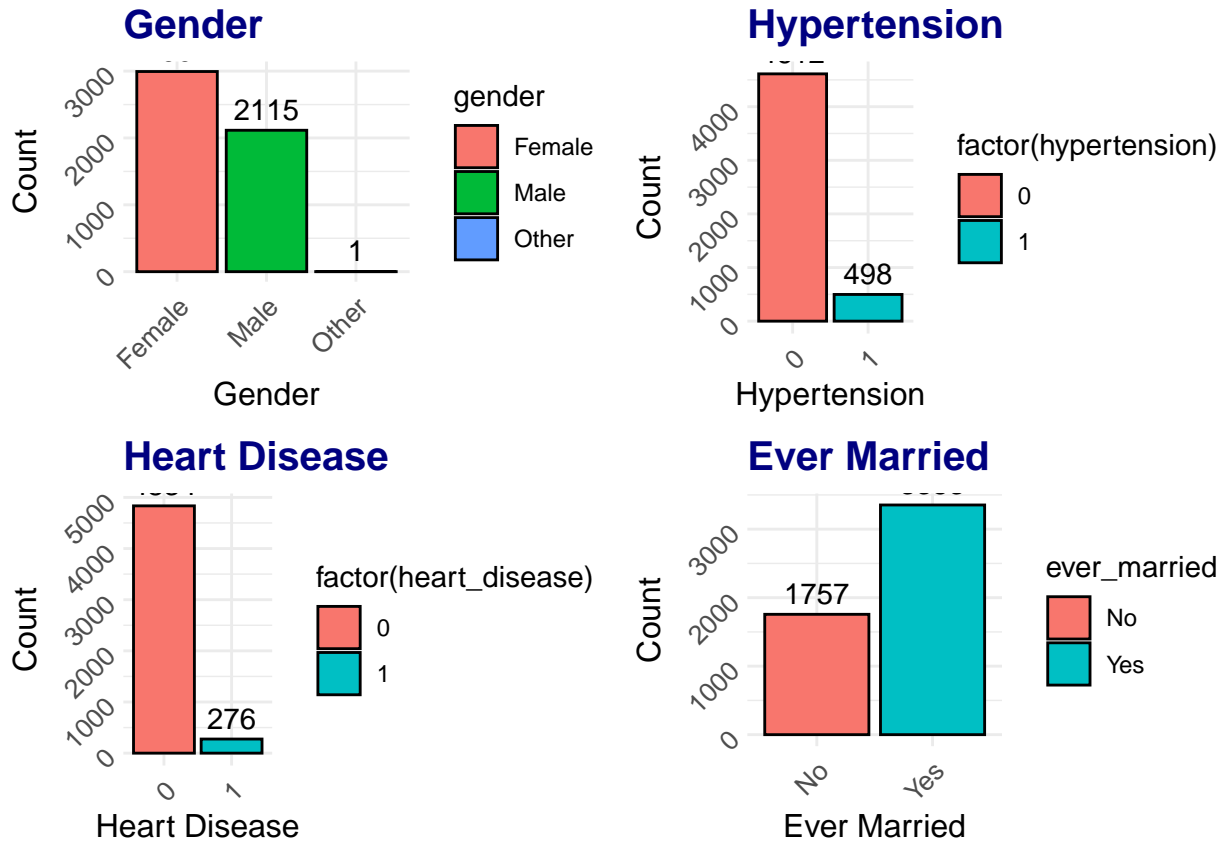
# Create a bar plot for smoking_status
p7 <- ggplot(data, aes(x = smoking_status)) +
  geom_bar(aes(fill = smoking_status), color = "black") +
  theme_minimal() +
  labs(title = "Smoking Status", x = "Smoking Status", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))

# Create a bar plot for stroke
p8 <- ggplot(data, aes(x = factor(stroke))) +
  geom_bar(aes(fill = factor(stroke)), color = "black") +
  theme_minimal() +
  labs(title = "Stroke", x = "Stroke", y = "Count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5) +
  theme(plot.title = element_text(size = 15, face = "bold", color = "navy"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10, angle = 45, hjust = 1))





```

```
# Combine the bar plots into one layout
grid.arrange(p1, p2, p3, p4, ncol = 2, nrow = 2)
```

```
## Warning: The dot-dot notation (`..count..`) was deprecated in ggplot2 3.4.0.
## i Please use `after_stat(count)` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```



```
grid.arrange(p5, p6, p7, p8, ncol = 2, nrow = 2)
```



```
# Load necessary libraries
library(ggplot2)
library(dplyr)

# Prepare the data
stroke_data <- data %>%
  group_by(stroke) %>%
  summarise(count = n()) %>%
  mutate(prop = count / sum(count) * 100,
         label = paste0(round(prop, 1), "%"))


# Create the pie chart with percentage labels
pie_chart <- ggplot(stroke_data, aes(x = "", y = prop, fill = factor(stroke))) +
  geom_bar(width = 1, stat = "identity") +
  coord_polar(theta = "y") +
  labs(title = "Pie Chart of Stroke Distribution", fill = "Stroke") +
  theme_minimal() +
  theme(axis.title.x = element_blank(),
        axis.title.y = element_blank(),
        panel.grid = element_blank(),
        axis.text.x = element_blank(),
        plot.title = element_text(size = 18, face = "bold"),
        legend.title = element_text(size = 14),
        legend.text = element_text(size = 12)) +
  geom_text(aes(label = label), position = position_stack(vjust = 0.5), size = 4)

# Display the pie chart
```



```
print(pie_chart)
```

Pie Chart of Stroke Distribution

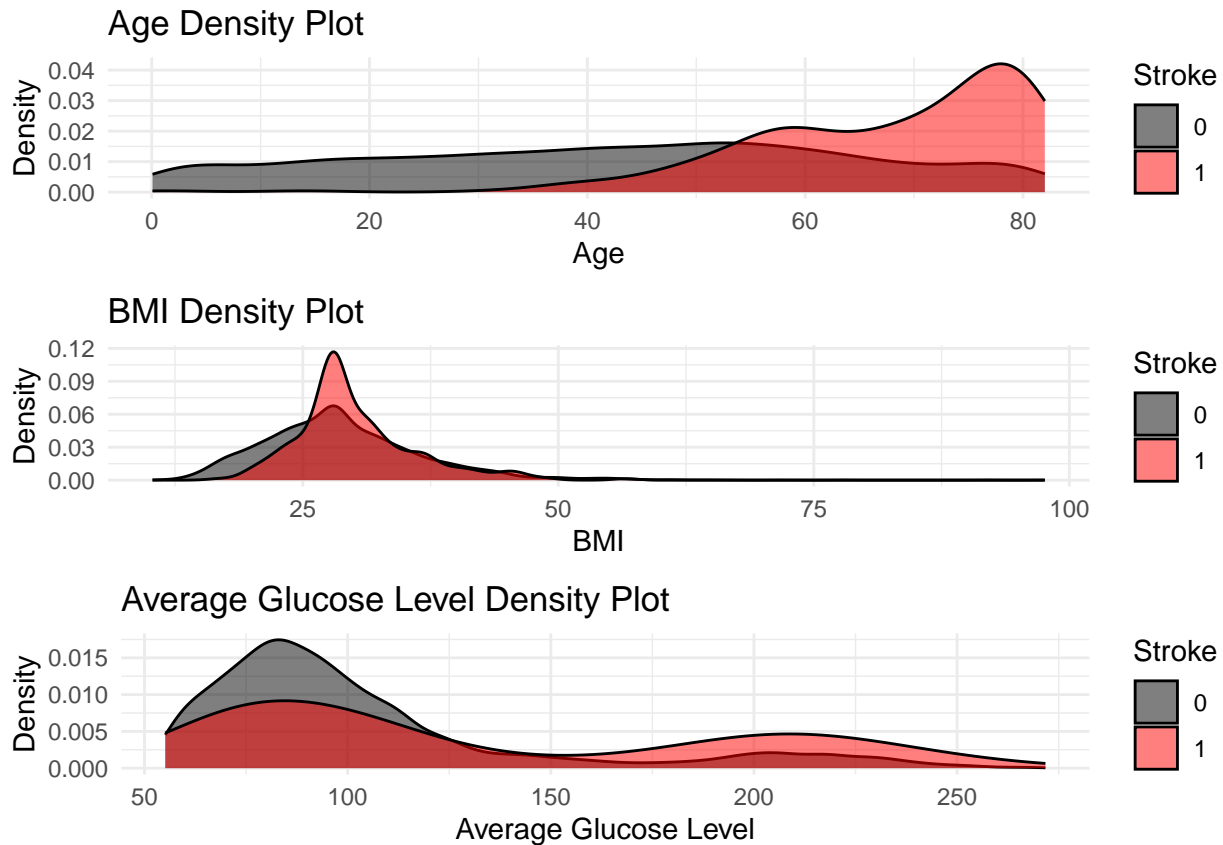


```
# Density plot for age
p1 = ggplot(data, aes(x = age, fill = factor(stroke))) +
  geom_density(alpha = 0.5) +
  theme_minimal() +
  labs(title = 'Age Density Plot', x = 'Age', y = 'Density', fill = 'Stroke') +
  scale_fill_manual(values = c("0" = "black", "1" = "red"))

# Density plot for bmi
p2 = ggplot(data, aes(x = bmi, fill = factor(stroke))) +
  geom_density(alpha = 0.5) +
  theme_minimal() +
  labs(title = 'BMI Density Plot', x = 'BMI', y = 'Density', fill = 'Stroke') +
  scale_fill_manual(values = c("0" = "black", "1" = "red"))

# Density plot for avg_glucose_level
p3 = ggplot(data, aes(x = avg_glucose_level, fill = factor(stroke))) +
  geom_density(alpha = 0.5) +
  theme_minimal() +
  labs(title = 'Average Glucose Level Density Plot', x = 'Average Glucose Level', y = 'Density', fill = 'Stroke') +
  scale_fill_manual(values = c("0" = "black", "1" = "red"))

grid.arrange(p1, p2, p3, ncol = 1, nrow = 3)
```



```
library(caret)
```

```
## Loading required package: lattice
```

```
data <- data %>%
  mutate(gender = as.factor(gender),
         ever_married = as.factor(ever_married),
         work_type = as.factor(work_type),
         Residence_type = as.factor(Residence_type),
         smoking_status = as.factor(smoking_status))
set.seed(123) # For reproducibility
train_index <- createDataPartition(data$stroke, p = 0.8, list = FALSE)
train_data <- data[train_index, ]
test_data <- data[-train_index, ]
# Check the dimensions of the split data
dim(train_data)
```

```
## [1] 4088  12
```

```
dim(test_data)
```

```
## [1] 1022  12
```

```
library(brms)
```

```
## Loading required package: Rcpp
```

```
## Loading 'brms' package (version 2.21.0). Useful instructions
## can be found by typing help('brms'). A more detailed introduction
## to the package is available through vignette('brms_overview').
```

```

##
## Attaching package: 'brms'

## The following object is masked from 'package:rstan':
##
##      loo

## The following object is masked from 'package:stats':
##
##      ar

priors <- c(set_prior("normal(0, 10)", class = "b"), # Prior for coefficients
            set_prior("student_t(3, 0, 10)", class = "Intercept")) # Prior for intercept
formula <- bf(stroke ~ gender + age + hypertension + heart_disease + ever_married +
              work_type + Residence_type + avg_glucose_level + bmi + smoking_status,
              family = bernoulli())

# Fit the model
fit <- brm(formula, data = train_data, prior = priors, iter = 2000, chains = 4, seed = 123)

## Compiling Stan program...
## Trying to compile a simple C file

## Running /Library/Frameworks/R.framework/Resources/bin/R CMD SHLIB foo.c
## using C compiler: 'Apple clang version 15.0.0 (clang-1500.3.9.4)'
## using SDK: 'MacOSX14.4.sdk'
## clang -arch arm64 -I"/Library/Frameworks/R.framework/Resources/include" -DNDEBUG -I"/Library/Frameworks/R.framework/Resources/include" -c foo.c
## In file included from <built-in>:1:
## In file included from /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/StanHeaders/include/stan/math/prim/mat/fun/eigen.h:1:
## In file included from /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/RcppEigen/include/Eigen/Core:1:
## In file included from /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/RcppEigen/include/Eigen/src/Core/Matrix.h:1:
## In file included from /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/RcppEigen/include/Eigen/src/Core/MatrixBase.h:1:
## In file included from /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/RcppEigen/include/Eigen/src/Core/MatrixBase.h:1:
## #include <cmath>
##      ~~~~~~
## 1 error generated.
## make: *** [foo.o] Error 1

## Start sampling

##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0.000112 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 1.12 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)

```

```

## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 3.768 seconds (Warm-up)
## Chain 1: 1.737 seconds (Sampling)
## Chain 1: 5.505 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 9.2e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.92 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 4.354 seconds (Warm-up)
## Chain 2: 1.77 seconds (Sampling)
## Chain 2: 6.124 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 9.4e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.94 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 4.305 seconds (Warm-up)

```

```

## Chain 3:          1.794 seconds (Sampling)
## Chain 3:          6.099 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 8.8e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.88 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 5.29 seconds (Warm-up)
## Chain 4:          1.802 seconds (Sampling)
## Chain 4:          7.092 seconds (Total)
## Chain 4:

```

```
summary(fit)
```

```

## Family: bernoulli
## Links: mu = logit
## Formula: stroke ~ gender + age + hypertension + heart_disease + ever_married + work_type + Residence
## Data: train_data (Number of observations: 4088)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##          total post-warmup draws = 4000
##
## Regression Coefficients:
##              Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS
## Intercept          -6.76      0.91   -8.78    -5.24 1.00     1755
## genderMale          -0.01      0.16   -0.33     0.31 1.00     6636
## genderOther         -5.01      7.02  -20.60     5.54 1.00     4798
## age                  0.08      0.01    0.06     0.09 1.00     3355
## hypertension         0.57      0.18    0.20     0.92 1.00     6428
## heart_disease        0.29      0.22   -0.15     0.69 1.00     5742
## ever_marriedYes     -0.13      0.27   -0.64     0.41 1.00     5615
## work_typeGovt_job   -1.10      0.99   -2.74     1.06 1.00     1246
## work_typeNever_worked -6.63      6.34  -21.67     2.31 1.00     3711
## work_typePrivate    -0.84      0.96   -2.45     1.33 1.00     1289
## work_typeSelfEmployed -1.30      0.99   -2.97     0.90 1.00     1352
## Residence_typeUrban   0.03      0.16   -0.27     0.34 1.00     5682
## avg_glucose_level     0.00      0.00    0.00     0.01 1.00     4522
## bmi                 -0.00      0.01   -0.03     0.02 1.00     5448
## smoking_statusneversmoked -0.18      0.20   -0.58     0.21 1.00     3913

```

```

## smoking_statussmokes          0.08      0.24     -0.40      0.56 1.00      3852
## smoking_statusUnknown        -0.16      0.24     -0.64      0.33 1.00      3759
##                               Tail_ESS
## Intercept                     1701
## genderMale                    3092
## genderOther                   2556
## age                           2743
## hypertension                  3100
## heart_disease                 2804
## ever_marriedYes              2944
## work_typeGovt_job             1551
## work_typeNever_worked        2587
## work_typePrivate              1487
## work_typeSelfMemployed        1524
## Residence_typeUrban           2814
## avg_glucose_level             2916
## bmi                           3130
## smoking_statusneversmoked     3185
## smoking_statussmokes          3058
## smoking_statusUnknown         2986
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

predictions <- predict(fit, newdata = test_data, type = "response")

# Convert probabilities to binary predictions
test_data$predicted_stroke <- ifelse(predictions[,1] > 0.5, 1, 0)

# Evaluate the model
conf_matrix <- table(test_data$stroke, test_data$predicted_stroke)
accuracy <- sum(diag(conf_matrix)) / sum(conf_matrix)
cat("Accuracy:", accuracy, "\n")

## Accuracy: 0.9452055

print(conf_matrix)

##
##      0   1
## 0 965   1
## 1  55   1

library(bayesplot)

## This is bayesplot version 1.11.1
## - Online documentation and vignettes at mc-stan.org/bayesplot
## - bayesplot theme set to bayesplot::theme_default()
##   * Does _not_ affect other ggplot2 plots
##   * See ?bayesplot_theme_set for details on theme setting
##
## Attaching package: 'bayesplot'

```

```
## The following object is masked from 'package:brms':
```

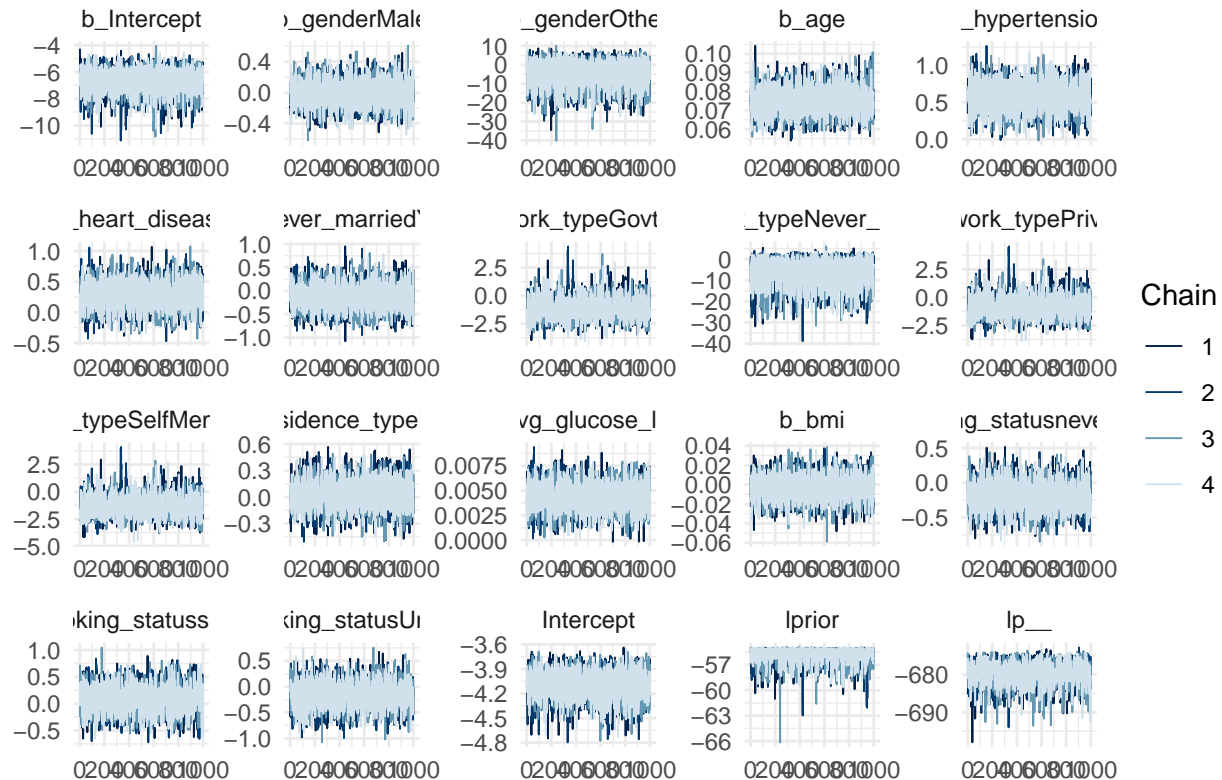
```
##
```

```
##      rhat
```

```
mcmc_trace(as.mcmc(fit)) + ggtitle("Trace Plots for Model Parameters") + theme_minimal()
```

```
## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
```

Trace Plots for Model Parameters



```
# Create individual plots for each parameter
```

```
g1 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_Intercept"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of the Intercept") +
  xlab("Intercept") +
  ylab("Density") +
  theme_minimal()
```

```
## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
```

```
g2 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_genderMale"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of genderMale") +
  xlab("genderMale") +
  ylab("Density") +
  theme_minimal()
```

```

theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
g3 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_genderOther"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of genderOther") +
  xlab("genderOther") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
# Continue for other parameters
g4 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_age"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of age") +
  xlab("age") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
g5 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_hypertension"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of hypertension1") +
  xlab("hypertension") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
g6 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_heart_disease"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of heart_disease1") +
  xlab("heart_disease") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
g7 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_ever_marriedYes"),
  prob = 0.95

```



```

) +
  ggtitle("Posterior Distribution of ever_marriedYes") +
  xlab("ever_marriedYes") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.

g8 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_work_typeGovt_job"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of work_typeGovt_job") +
  xlab("work_typeGovt_job") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.

g9 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_work_typeNever_worked"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of work_typeNever_worked") +
  xlab("work_typeNever_worked") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.

g10 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_work_typePrivate"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of work_typePrivate") +
  xlab("work_typePrivate") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.

g11 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_work_typeSelfEmployed"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of work_typeSelfEmployed") +
  xlab("work_typeSelfEmployed") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.

```

```
g12 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_Residence_typeUrban"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of Residence_typeUrban") +
  xlab("Residence_typeUrban") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
```

```
g13 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_avg_glucose_level"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of avg_glucose_level") +
  xlab("avg_glucose_level") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
```

```
g14 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_bmi"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of bmi") +
  xlab("bmi") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
```

```
g15 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_smoking_statusneversmoked"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of smoking_statusneversmoked") +
  xlab("smoking_statusneversmoked") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
```

```
g16 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_smoking_statussmokes"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of smoking_statussmokes") +
  xlab("smoking_statussmokes") +
  ylab("Density") +
```

```

theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
g17 <- mcmc_areas(
  as.mcmc(fit),
  pars = c("b_smoking_statusUnknown"),
  prob = 0.95
) +
  ggtitle("Posterior Distribution of smoking_statusUnknown") +
  xlab("smoking_statusUnknown") +
  ylab("Density") +
  theme_minimal()

## Warning: as.mcmc.brmsfit is deprecated and will eventually be removed.
# Combine all plots into a grid
grid.arrange(g1, g2, g3, g4, g5, g6, g7, g8, g9, g10, g11, g12, g13, g14, g15, g16, g17, ncol = 3)

```

```

predictions <- predict(fit, newdata = test_data, type = "response")

# Convert probabilities to binary predictions
test_data$predicted_stroke <- ifelse(predictions[,1] > 0.3, 1, 0)

# Evaluate the model
conf_matrix <- table(test_data$stroke, test_data$predicted_stroke)
accuracy <- sum(diag(conf_matrix)) / sum(conf_matrix)
cat("Accuracy:", accuracy, "\n")

```

```

## Accuracy: 0.9363992
print(conf_matrix)

##
##      0    1
## 0 952  14
## 1  51   5

# Load necessary libraries
library(brms)
library(dplyr)

# Assuming 'fit' is your fitted model object
# and 'train_data' is your training dataset

# Generate posterior predictions for the training data
y_rep_train <- posterior_predict(fit, newdata = train_data)

# Calculate the mean predicted probability across the posterior samples
predicted_probs_train <- apply(y_rep_train, 2, mean)

# Set a threshold for stroke prediction (e.g., 0.5)
threshold <- 0.25

# Classify the individual based on the predicted probability
predicted_class_train <- ifelse(predicted_probs_train > threshold, 1, 0)

# Create a confusion matrix
confusion_matrix_train <- table(train_data$stroke, predicted_class_train)
print(confusion_matrix_train)

##      predicted_class_train
##      0      1
## 0 3812   83
## 1  166   27

# Calculate training accuracy
training_accuracy <- sum(diag(confusion_matrix_train)) / sum(confusion_matrix_train)
print(paste("Training Accuracy:", training_accuracy))

## [1] "Training Accuracy: 0.939090019569472"

# Create a data frame for the individual
individual_data <- data.frame(
  gender = "Male",
  age = 68,
  hypertension = 0,
  heart_disease = 1,
  ever_married = "Yes",
  work_type = "Private",
  Residence_type = "Urban",
  avg_glucose_level = 229,
  bmi = 36,
  smoking_status = "formerly smoked"
)

```

```

# Ensure the levels of categorical variables match the training data
individual_data$gender <- factor(individual_data$gender, levels = levels(train_data$gender))
individual_data$ever_married <- factor(individual_data$ever_married, levels = levels(train_data$ever_married))
individual_data$work_type <- factor(individual_data$work_type, levels = levels(train_data$work_type))
individual_data$Residence_type <- factor(individual_data$Residence_type, levels = levels(train_data$Residence_type))
individual_data$smoking_status <- factor(individual_data$smoking_status, levels = levels(train_data$smoking_status))

# Generate posterior predictions for the individual
predicted_probs <- posterior_predict(fit, newdata = individual_data)

# Calculate the mean predicted probability across the posterior samples
predicted_prob <- mean(predicted_probs)

# Print the predicted probability
print(predicted_prob)

```

```
## [1] 0.2015
```

```

# Define a threshold for classification (e.g., 0.5)
threshold <- 0.5

# Classify the individual based on the predicted probability
predicted_class <- ifelse(predicted_prob > threshold, 1, 0)

# Print the predicted class
if(predicted_class == 1) {
  cat("The model predicts that the individual is at risk of having a stroke.\n")
} else {
  cat("The model predicts that the individual is not at risk of having a stroke.\n")
}

```

```
## The model predicts that the individual is not at risk of having a stroke.
```

```

# Load necessary libraries
library(brms)
library(dplyr)

# Assuming 'fit' is your fitted model object
# and 'test_data' is your test dataset

# Generate posterior predictions for the test data
y_rep_test <- posterior_predict(fit, newdata = test_data)

# Calculate the mean predicted probability across the posterior samples
predicted_probs <- apply(y_rep_test, 2, mean)

# Define a list of thresholds
thresholds <- seq(0.1, 0.9, by = 0.01)

# Initialize a vector to store accuracy results
accuracies <- numeric(length(thresholds))

# Loop through each threshold and calculate accuracy
for (i in seq_along(thresholds)) {
  threshold <- thresholds[i]

```

```

predicted_class <- ifelse(predicted_probs > threshold, 1, 0)
confusion_matrix <- table(test_data$stroke, predicted_class)

# Calculate accuracy
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
accuracies[i] <- accuracy
}

# Combine thresholds and accuracies into a data frame
results <- data.frame(Threshold = thresholds, Accuracy = accuracies)

# Print the results
print(results)

```

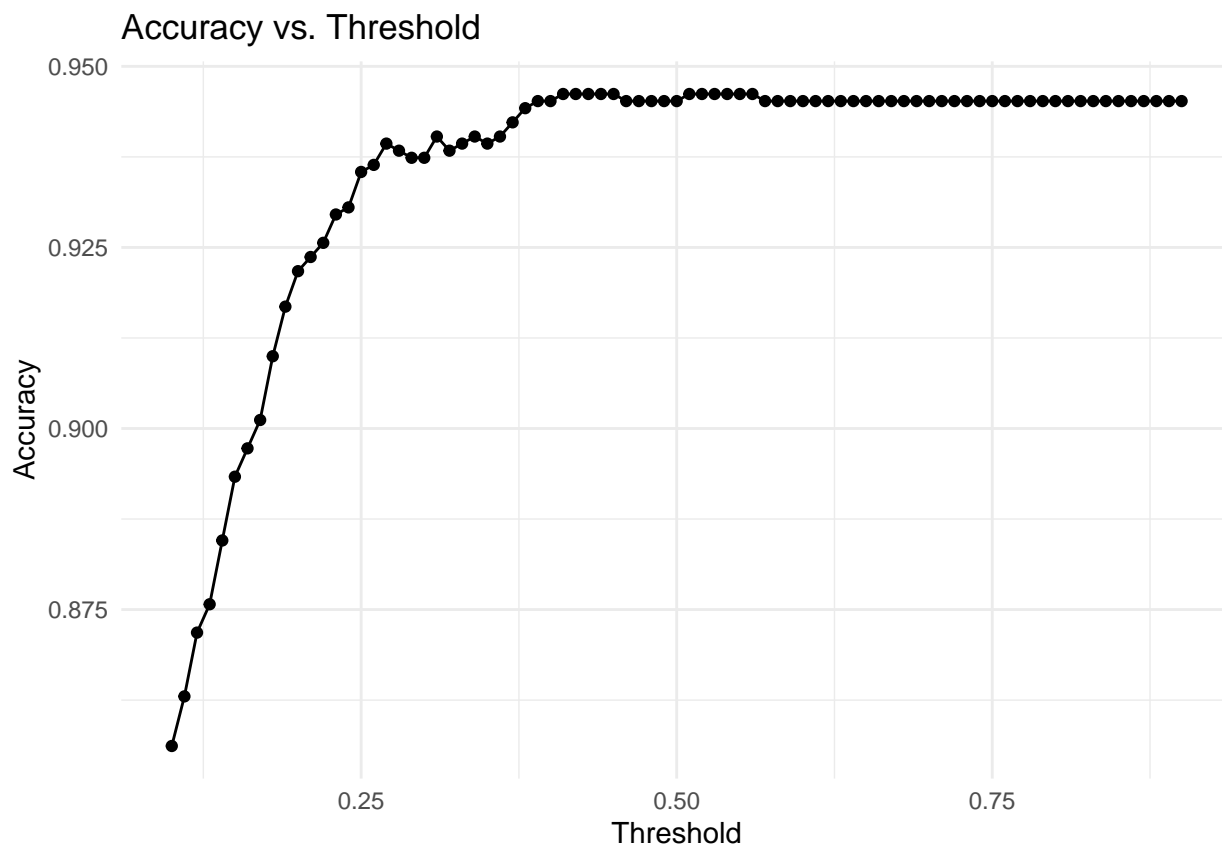
```

##      Threshold  Accuracy
## 1         0.10 0.8561644
## 2         0.11 0.8630137
## 3         0.12 0.8718200
## 4         0.13 0.8757339
## 5         0.14 0.8845401
## 6         0.15 0.8933464
## 7         0.16 0.8972603
## 8         0.17 0.9011742
## 9         0.18 0.9099804
## 10        0.19 0.9168297
## 11        0.20 0.9217221
## 12        0.21 0.9236791
## 13        0.22 0.9256360
## 14        0.23 0.9295499
## 15        0.24 0.9305284
## 16        0.25 0.9354207
## 17        0.26 0.9363992
## 18        0.27 0.9393346
## 19        0.28 0.9383562
## 20        0.29 0.9373777
## 21        0.30 0.9373777
## 22        0.31 0.9403131
## 23        0.32 0.9383562
## 24        0.33 0.9393346
## 25        0.34 0.9403131
## 26        0.35 0.9393346
## 27        0.36 0.9403131
## 28        0.37 0.9422701
## 29        0.38 0.9442270
## 30        0.39 0.9452055
## 31        0.40 0.9452055
## 32        0.41 0.9461840
## 33        0.42 0.9461840
## 34        0.43 0.9461840
## 35        0.44 0.9461840
## 36        0.45 0.9461840
## 37        0.46 0.9452055
## 38        0.47 0.9452055
## 39        0.48 0.9452055

```

```
## 40      0.49 0.9452055
## 41      0.50 0.9452055
## 42      0.51 0.9461840
## 43      0.52 0.9461840
## 44      0.53 0.9461840
## 45      0.54 0.9461840
## 46      0.55 0.9461840
## 47      0.56 0.9461840
## 48      0.57 0.9452055
## 49      0.58 0.9452055
## 50      0.59 0.9452055
## 51      0.60 0.9452055
## 52      0.61 0.9452055
## 53      0.62 0.9452055
## 54      0.63 0.9452055
## 55      0.64 0.9452055
## 56      0.65 0.9452055
## 57      0.66 0.9452055
## 58      0.67 0.9452055
## 59      0.68 0.9452055
## 60      0.69 0.9452055
## 61      0.70 0.9452055
## 62      0.71 0.9452055
## 63      0.72 0.9452055
## 64      0.73 0.9452055
## 65      0.74 0.9452055
## 66      0.75 0.9452055
## 67      0.76 0.9452055
## 68      0.77 0.9452055
## 69      0.78 0.9452055
## 70      0.79 0.9452055
## 71      0.80 0.9452055
## 72      0.81 0.9452055
## 73      0.82 0.9452055
## 74      0.83 0.9452055
## 75      0.84 0.9452055
## 76      0.85 0.9452055
## 77      0.86 0.9452055
## 78      0.87 0.9452055
## 79      0.88 0.9452055
## 80      0.89 0.9452055
## 81      0.90 0.9452055
```

```
# Optionally, plot the results
library(ggplot2)
ggplot(results, aes(x = Threshold, y = Accuracy)) +
  geom_line() +
  geom_point() +
  ggtitle("Accuracy vs. Threshold") +
  xlab("Threshold") +
  ylab("Accuracy") +
  theme_minimal()
```



```
# Load necessary libraries
```

```
library(brms)
```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
# Assuming 'fit' is your fitted model object
```

```
# and 'test_data' is your test dataset
```

```
# Generate posterior predictions for the test data
```

```
y_rep_test <- posterior_predict(fit, newdata = test_data)
```

```
# Calculate the mean predicted probability across the posterior samples
```

```
predicted_probs <- apply(y_rep_test, 2, mean)
```

```
# Calculate the ROC curve and AUC
```

```
roc_obj <- roc(test_data$stroke, predicted_probs)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```



```
# Print the AUC value
auc_value <- auc(roc_obj)
print(paste("ROC-AUC:", auc_value))

## [1] "ROC-AUC: 0.841162008281574"

# Plot the ROC curve
plot(roc_obj, main = paste("ROC Curve (AUC =", round(auc_value, 2), ")"))
```

